

Hands-on Lab : Create Tables and Load Data in PostgreSQL using pgAdmin

Estimated time needed: 20 minutes

In this lab, you will learn how to create tables and load data in the PostgreSQL database service using the pgAdmin graphical user interface (GUI) tool. The pgAdmin GUI provides an alternative to the command line for interacting with a PostgreSQL database using a graphical interface. This provides a number of key features for interacting with a PostgreSQL database in an easy to use format.

Software Used in this Lab

In this lab, you will use [PostgreSQL Database](#). PostgreSQL is a Relational Database Management System (RDBMS) designed to efficiently store, manipulate, and retrieve data.



To complete this lab you will utilize the PostgreSQL relational database service available as part of IBM Skills Network Labs (SN Labs) Cloud IDE. SN Labs is a virtual lab environment used in this course.

Database Used in this Lab

Books database has been used in this lab.

The following diagram shows the structure of the myauthors table from the Books database:

myauthors	
author_id	int
first_name	varchar(100)
middle_name	varchar(50)
last_name	varchar(100)

Objectives

After completing this lab, you will be able to use pgAdmin with PostgreSQL to:

- Create databases and tables in a PostgreSQL instance
- Load data into tables manually using the pgAdmin GUI
- Load data into tables from a text/script file

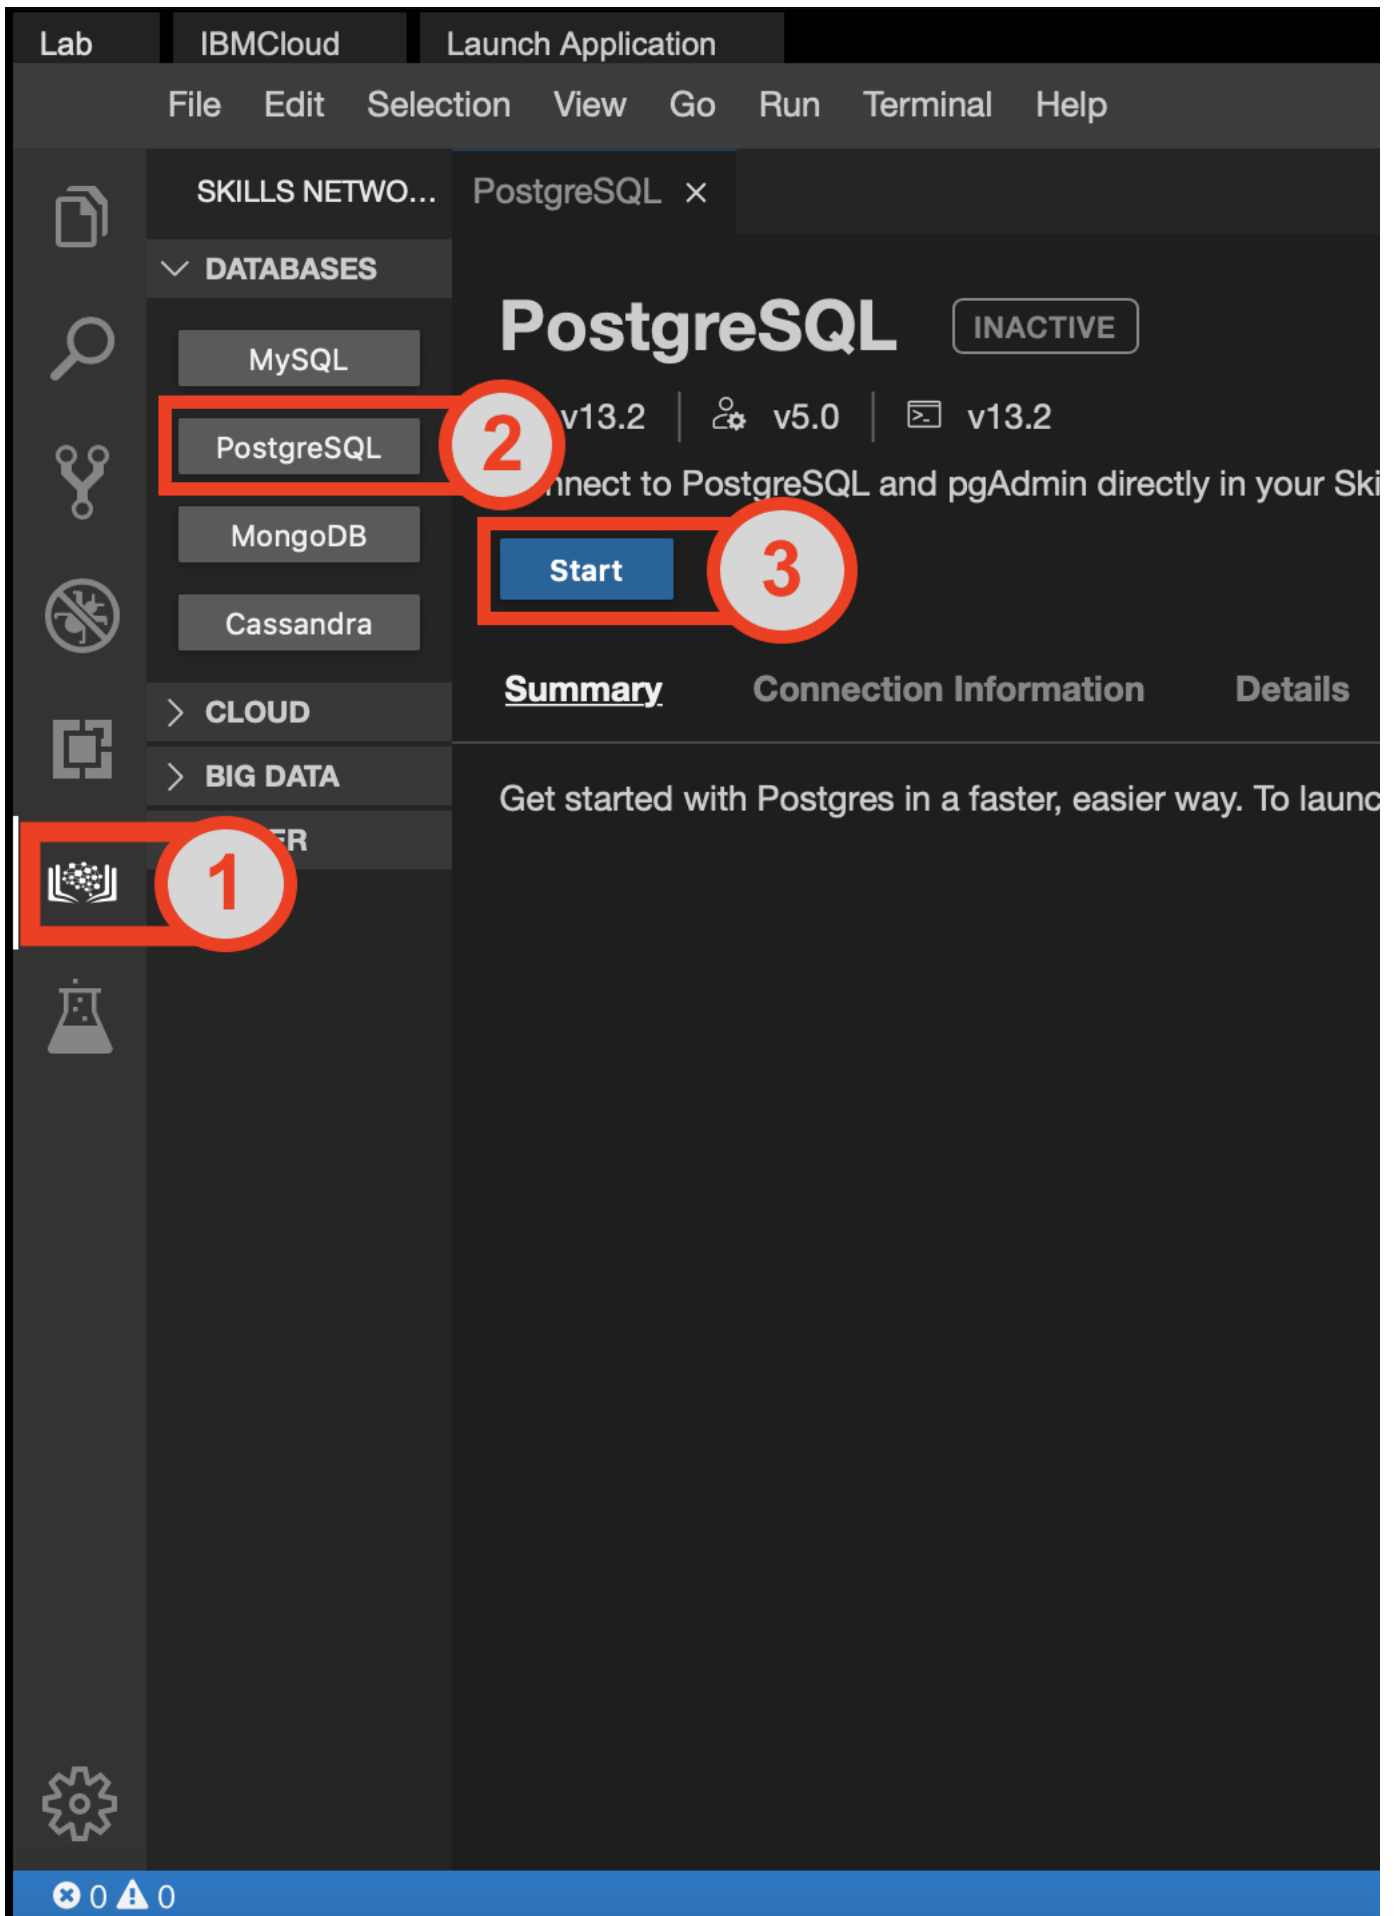
Lab Structure

In this lab, you will complete several tasks in which you will learn how to create tables and load data in the PostgreSQL database service using the pgAdmin graphical user interface (GUI) tool.


Task A: Create a database

First, to create a database on a PostgreSQL server instance, you'll first want to actually launch a PostgreSQL server instance on Cloud IDE and open up the pgAdmin Graphical User Interface.

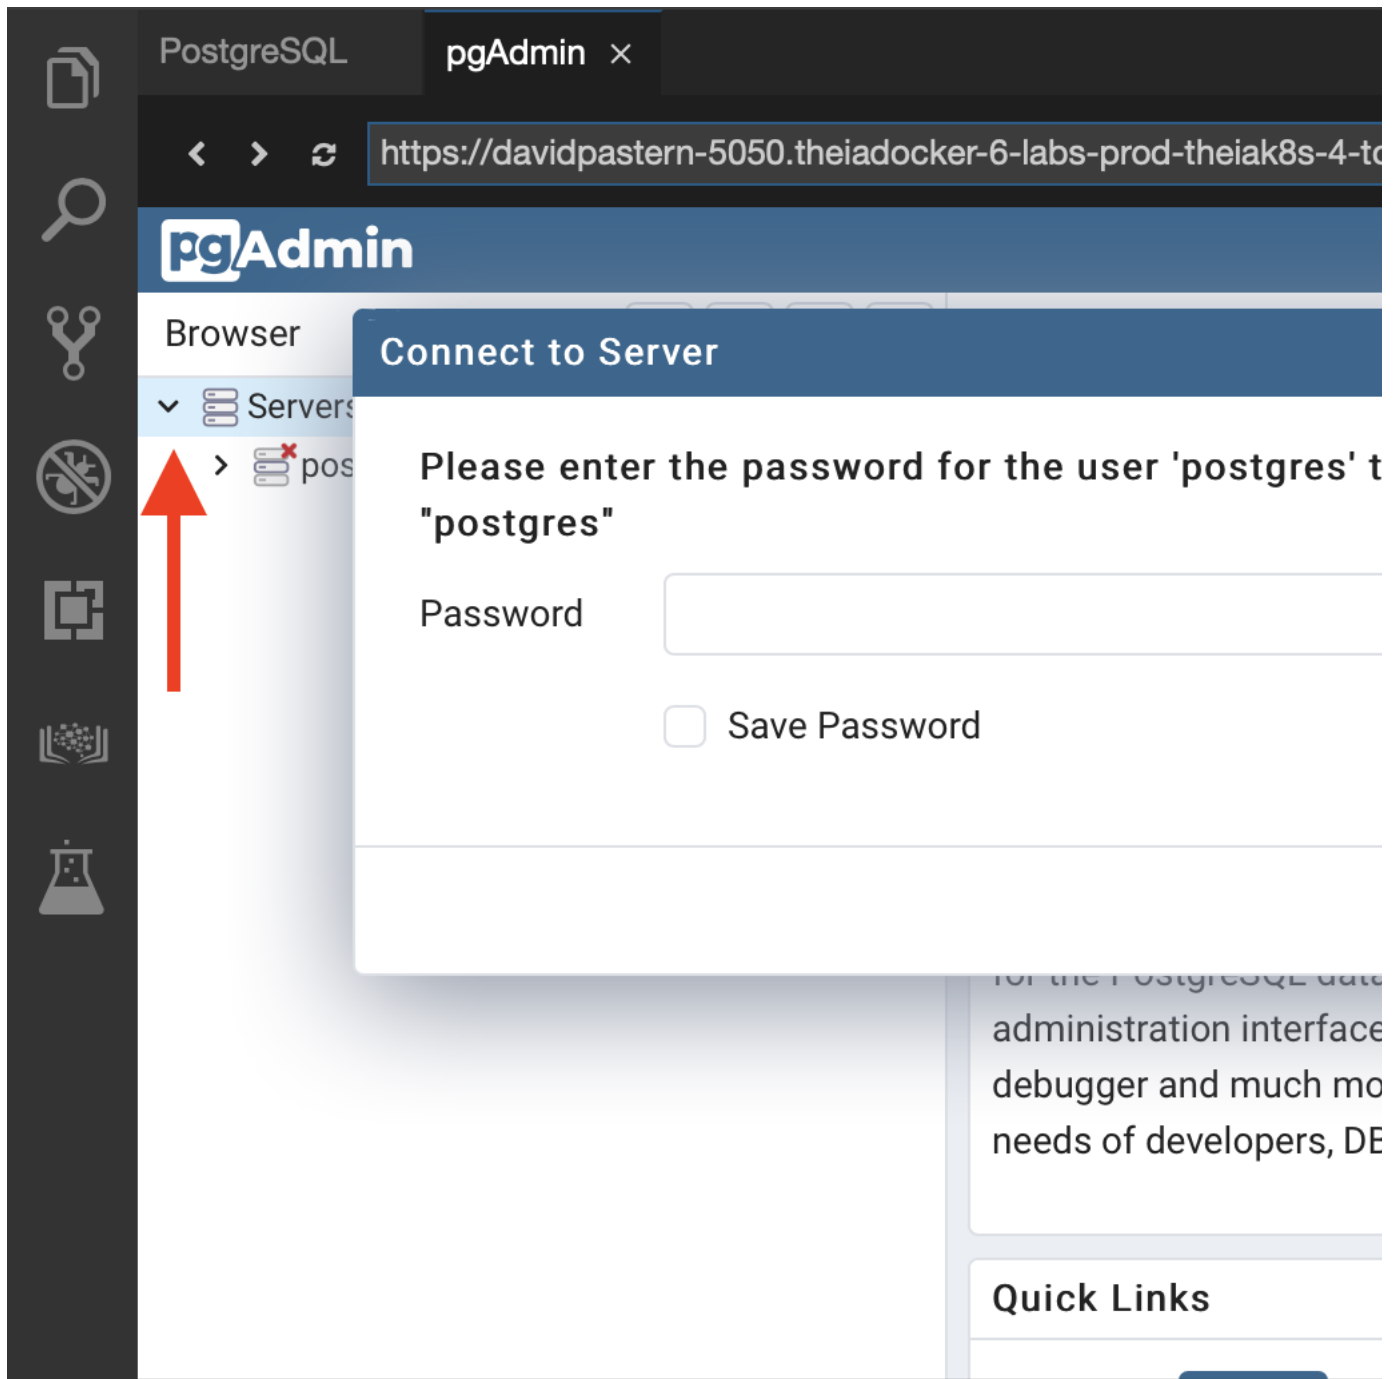
1. Click on the Skills Network extension button on the left side of the window.
2. Open the "DATABASES" drop down menu and click on "PostgreSQL"
3. Click on the "Start" button. PostgreSQL may take a few moments to start.



4. Next, open the pgAdmin Graphical User Interface by clicking the “pgAdmin” button in the Cloud IDE interface.



5. Once the pgAdmin GUI opens, click on the Servers tab on the left side of the page. You will be prompted to enter a password.



6. To retrieve your password, click on the “PostgreSQL” tab near the top of the interface.

7. Click on the Copy icon to the left of your password to copy the session password onto your clipboard.

PostgreSQL x pgAdmin

PostgreSQL

ACTIVE

v13.2 | v5.0 | v13.2

Connect to PostgreSQL and pgAdmin directly in your Skills Network Labs environment

Stop

Summary Connection Information Details

Your database and pgAdmin server are now ready to use and available with the terminal. For more information on how to navigate PostgreSQL, please check out the Details section.

Username:

Password:

You can manage PostgreSQL via:

pgAdmin

Or to interact with the database in the terminal, select one of these options:

PostgreSQL CLI New Terminal

8. Navigate back to the “pgAdmin” tab and paste in your password, then click OK

9. You will then be able to access the pgAdmin GUI tool.



Welcome



pgAdmin
Management

Feature rich | Maximizing

pgAdmin is an Open Source administration tool designed to answer the needs of PostgreSQL users.

Quick Links

Getting Started



PostgreSQL Documentation

10. In the tree-view, expand **Servers** > **postgres** > **Databases**. If prompted, enter your PostgreSQL service session password. Right-click on **Databases** and go to **Create** > **Database**. In the **Database** box, type **Books** as the name for your new database, and then click **Save**. Proceed to Task B.

The screenshot shows the pgAdmin interface. The top navigation bar includes the pgAdmin logo and menus for File, Object, Tools, and Help. Below this is a toolbar with icons for a browser, table, filter, and search. The main area is divided into a left tree view and a right pane. The tree view shows a hierarchy: Servers (1) > postgres > Databases (1). The 'Databases (1)' node is selected and right-clicked, opening a context menu with options: Create, Refresh..., and Database... (highlighted with a mouse cursor). The right pane shows 'Server sessions' with a table of 7 sessions and 'Tuples in' with a table of 1 tuple.

pgAdmin File ▾ Object ▾ Tools ▾ Help ▾

Browser 1

2

3

▼ Servers (1)

▼ postgres

▼ Databases (1)

▼ postgres

► Casts

► Catalogs

► Event Triggers

► Extensions

► Foreign Data Wrappers

► Languages

► Publications

► Schemas

► Subscriptions

► Login/Group Roles

► Tablespaces

Create > Database...

Refresh...

Server sessions

7
4
3
2
1
0

Tuples in

1

Create - Database

General

Definition

Security

Parameters


Advanced

SQL

Database

Books

Owner

 postgres

Comment

i

?

✕ Cancel

Task B: Create tables

Now that you have your PostgreSQL service active and have created the **Books** database using pgAdmin, let's go ahead and create a few tables to populate the database and store the data that we wish to eventually upload into it.

1. In the tree-view, expand **Books** > **Schemas** > **public**. Right-click on **Tables** and go to **Create** > **Table**.

Browser



Dashboard

Properties

SQL

▾ Servers (1)

▾ postgres

▾ Databases (2)

▾ Books 1

> Casts

> Catalogs

> Event Triggers

> Extensions

> Foreign Data Wrappers

> Languages

> Publications

▾ Schemas (1) 2

▾ public 3

> Collations

> Domains

> FTS Configurations

> FTS Dictionaries

> FTS Parsers

> FTS Templates

> Foreign Tables

> Functions

> Materialized Views

> Procedures

> 1..3 Sequences

4 > Tables

> Trigger

> Types

> Views

> Subscriptions

> postgres

> Login/Group Roles

> Tablespaces

Database sessions

1

0

Tuples in

1

0

Server activity

Sessions

Locks

Prepared

			PID	User

Create



Table...

Refresh...

Grant Wizard...

Search Objects...

Query Tool

2. On the **General** tab, in the **Name** box, type **myauthors** as name of the table. Don't click Save, proceed to the next step.

Create - Table

General

Columns

Advanced

Constraints

Partitions

Parameters

Se

Name

myauthors

Owner

postgres

Schema

public

Tablespace

Select an item...

Partitioned table?

No

Comment

i

?

Cancel

3. Switch to tab **Columns** and click the **Add new row** button four times to add **4** column placeholders. Don't click Save, proceed to the next step.

Create - Table

General

Columns

Advanced

Constraints

Partitions









Parameters

Se

Inherited from table(s)

Select to inherit from...

Columns

		Name ▲	Data type	Length/Precision	Scale
		<input type="text"/>	Select an item... ▼		
		<input type="text"/>	Select an item... ▼		
		<input type="text"/>	Select an item... ▼		
		<input type="text"/>	Select an item... ▼		

i

?

✕ Cancel

4. Enter the **myauthors** table definition structure information as shown in the image below in the highlighted boxes. Then click **Save**. Proceed to Task C.

Create - Table

GeneralColumnsAdvancedConstraintsPartitionsParametersSe

Inherited from table(s)Select to inherit from...

Columns

		Name	Data type	Length/Precision	Scale
		author_id	integer		
		first_name	character varying	100	
		middle_name	character varying	50	
		last_name	character varying	100	

i?

✕ Cancel

Task C: Load data into tables manually using the pgAdmin GUI

Great! You now have a database and have created tables within it. With the pgAdmin GUI, you can insert values into the tables manually. This is useful if you have a few new entries you wish to add to the database. Let's see how to do it.

1. In the tree-view, expand **Tables**. Right-click on **myauthors** and go to **View/Edit Data > All Rows**.

Servers (1)

- postgres
 - Databases (2)
 - Books
 - Casts
 - Catalogs
 - Event Triggers
 - Extensions
 - Foreign Data Wrap
 - Languages
 - Publications
 - Schemas (1)
 - public
 - Collations
 - Domains
 - FTS Config
 - FTS Diction
 - FTS Parser
 - FTS Templ
 - Foreign Tal
 - Functions
 - Materialize
 - Procedures
 - Sequences
 - 1 Tables (1)
 - myauthors
 - Columns
 - Constraints (1)
 - Indexes
 - RLS Policies
 - Rules
 - Triggers

Create

Refresh...

Count Rows

Delete/Drop

Drop Cascade

Reset Statistics

Import/Export...

Maintenance...

Scripts

Truncate

Backup...

Restore...

View/Edit Data

Search Objects...

Query Tool

Properties...

Type

Primary Key

All Rows

First 100 Row

Last 100 Row

Filtered Rows

2. You will insert 2 rows of data into the **myauthors** table. In the lower **Data Output** pane, enter **myauthors** table data information for 2 rows as shown in the highlighted boxes in the image below. Then click the **Save Data Changes** button. Proceed to Task D.

The screenshot shows the pgAdmin interface. At the top, there are tabs for Dashboard, Properties, SQL, Statistics, Dependencies, and Dependents. Below these is a toolbar with various icons. One icon, representing saving data changes, is highlighted with a red box. A callout bubble points to this icon with the text "Save Data Changes icon". Below the toolbar is the Query Editor, which contains the following SQL query:

```
1 SELECT * FROM pub
2 ORDER BY author_id ,
```

At the bottom of the interface is the Data Output pane, which shows the results of the query. It has tabs for Data Output, Explain, Messages, and Notifications. The Data Output tab is selected, and it displays a table with the following columns: author_id [PK] integer, first_name character varying (100), middle_name character varying (50), and last_name character varying (50). The table contains two rows of data:

	author_id [PK] integer	first_name character varying (100)	middle_name character varying (50)	last_name character varying (50)
1	1	Merrit	[null]	Eric
2	2	Linda	[null]	Mul

The first two columns of the first two rows are highlighted with red boxes.

Task D: Load data into tables using a text/script file

In the previous task, you entered some data entries into a table manually with pgAdmin. While this method can be useful for small additions, if you wish to upload large amounts of data at once, that process becomes far too tedious. An alternative is to load data into tables from a text or script file containing the data you wish to enter. Let's take a look at how to do this.

1. Finally, you will import the remainder of the **myauthors** table data from a csv text file. Download the csv file below to your local computer:

- [myauthors.csv](#)

2. In the tree-view, right-click on **myauthors** and go to **Import/Export**.

pgAdmin File Object Tools Help

Browser Dashboard Pr

- Servers (1)
 - postgres
 - Databases (2)
 - Books
 - Casts
 - Catalogs
 - Event Triggers
 - Extensions
 - Foreign Data W
 - Languages
 - Publications
 - Schemas (1)
 - public
 - Collatio
 - Domain
 - FTS Co
 - FTS Dic
 - FTS Par
 - FTS Ter
 - Foreign
 - Funcio
 - Materia
 - Procedu
 - 1..3 Sequen
 - 1 Tables
 - 2 mya
 - Columns
 - Constraints (1)
 - Indexes
 - RLS Policies
 - Rules
 - Triggers

Query Editor

```
1 SELECT *
```

2

Create

- Refresh...
- Count Rows
- Delete/Drop
- Drop Cascade
- Reset Statistics
- Import/Export...
- Maintenance...
- Scripts
- Truncate
- Backup...
- Restore...
- View/Edit Data
- Search Objects...
- Query Tool
- Properties...

3. Follow the instructions below to import:

- Make sure Import/Export is set to **Import**, Format = **csv** and Header = **Yes**. Then click on the **Select file** button by the Filename box.

Import/Export data - table 'myauthors'

Options

Columns

Import/Export

Import

1

File Info

Filename

Format

csv

2

Encoding

Select an item...

Miscellaneous

OID

No

Header

Yes

3

Delimiter

Select from list...

Specifies the character that separates columns with a file. The default is a tab character in text format, a character must be a single one-byte character. This option is not available in binary format.

- Click the **Upload File** button.

Select file



/var/lib/pgadmin/




Name	Size
 sessions	4.0 kB
 storage	4.0 kB




Show hidden files and folders? ☐


- Double-click on the drop files area and load the **myauthors.csv** you downloaded earlier from your local computer storage.

Select file





/var/lib/pgadmin/



Double click on this space

Drop files here to upload. The file size limit (per file) is

Show hidden files and folders?☐

- When the upload is complete, close the drop files area clicking the X button.

Select file



/var/lib/pgadmin/



26.6 KB



myauthors.csv

100%

Drop files here to upload. The file size limit (per file) is

Show hidden files and folders? ☐




- Select the uploaded **myauthors.csv** file from the list and click the **Select** button.

Select file



/var/lib/pgadmin/myauthors.csv



Name	Size
 myauthors.csv	26.0 kB
 sessions	4.0 kB
 storage	4.0 kB

Show hidden files and folders? ☐

- Click **OK** and notification of import success should appear.

Import/Export data - table 'myauthors'

Options

Columns

Import/Export

Import

File Info

Filename

/var/lib/pgadmin/myauthors.csv

Format

csv

Encoding

Select an item...

Miscellaneous

OID

No

Header

Yes

Delimiter

Select from list...

Specifies the character that separates columns with a file. The default is a tab character in text format, a comma in binary format. This option is not applicable in binary format.

Import - Copying table data



Copying table data 'public.myauthors' on database 'Books' and server (postgres:5432)

Mon Mar 22 2021 02:26:40 GMT-0600 (Mountain Daylight Time)



0.02 seconds

More details...

Stop Process



Successfully completed.

4. Repeat Task C Step 1 to check that the newly imported data rows appear along with your previously inserted 2 rows.

Query Editor
Query History

```

1 SELECT * FROM public.myauthors
2 ORDER BY author_id ASC

```

Data Output
Explain
Messages
Notifications

	author_id [PK] integer		first_name character varying (100)		middle_name character varying (50)	
1		1	Merrit		[null]	
2		2	Linda		[null]	
3		3	Alecos		[null]	
4		4	Paul		C.van	
5		5	David		[null]	
6		6	Richard		[null]	
7		7	Yuval		Noah	
8		8	Paul		[null]	
9		9	David		[null]	
10		10	John		Paul	
11		11	Andrew		[null]	
12		12	Melanie		[null]	
13		13	Neal		[null]	
14		14	Nir		[null]	
15		15	Tim		[null]	
16		16	Mike		[null]	
17		17	Brian		P.	
18		18	Jean-Philippe		[null]	
19		19	Lance		[null]	
20		20	Richard		C.	
21		21	William		L.	

22	22	Magnus	Lie
23	23	Mike	[null]
24	24	Norman	[null]
25	25	John	E.
26	26	S.	[null]

As you can see, the data contained in the `csv` file was successfully uploaded into the table and you did not have to manually input hundreds of entries.

Conclusion

Congratulations! You have completed this lab, and you are ready for the next topic.

Author

- [Sandip Saha Joy](#)

Other Contributors

- [David Pasternak](#)

Changelog

Date	Version	Changed by	Change Description
2021-03-15	1.0	Sandip Saha Joy	Created initial version
2021-10-18	1.1	David Pasternak	Updated lab instructions

© IBM Corporation 2020. All rights reserved.